



Greedy algorithms 1

Scheduling

CS240

Spring 2024

Rui Fan

Greedy algorithms

- Make the best choice at the moment.
 - No planning ahead. “Short-sighted”.
- Once choice made, it’s fixed.
 - No take-backs.
- **Cons** Doesn’t always find optimal answer.
- **Pros** Simple and fast. Sometimes optimal.





Overview

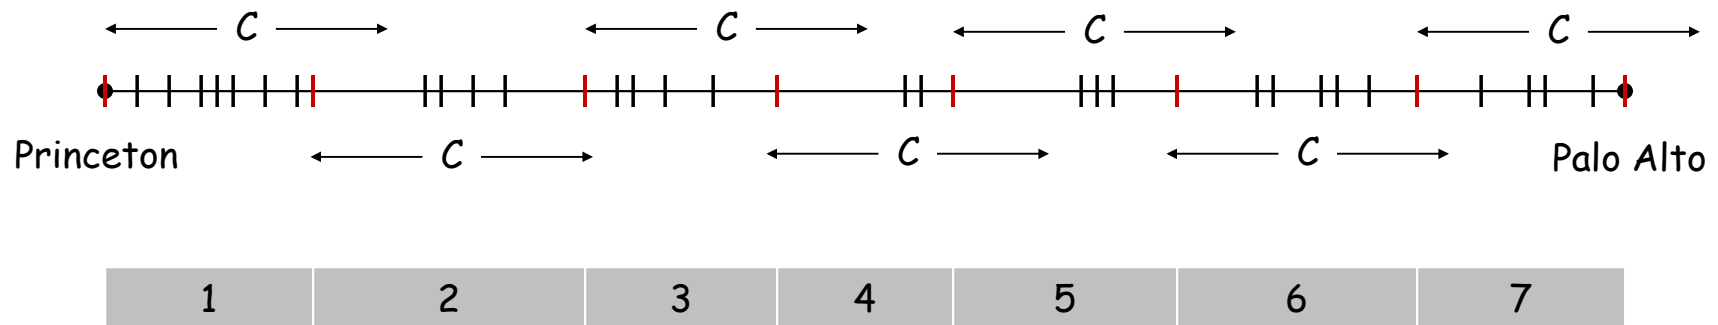
- Selecting breakpoints
- Coin change
- Interval scheduling
- Interval coloring
- Scheduling to minimizing lateness

Selecting Breakpoints

Selecting breakpoints.

- Road trip from Princeton to Palo Alto along fixed route.
- Refueling stations at certain points along the way.
- Fuel capacity = C .
- Goal: makes as few refueling stops as possible.

Greedy algorithm. Go as far as you can before refueling.



Selecting Breakpoints: Greedy Algorithm

Truck driver's algorithm.

```
Sort breakpoints so that:  $0 = b_0 < b_1 < b_2 < \dots < b_n = L$ 
```

```
 $S \leftarrow \{0\}$   $\leftarrow$  breakpoints selected
```

```
 $x \leftarrow 0$   $\leftarrow$  current location
```

```
while ( $x \neq b_n$ )
```

```
    let  $p$  be largest integer such that  $b_p \leq x + C$ 
```

```
    if ( $b_p = x$ )
```

```
        return "no solution"
```

```
     $x \leftarrow b_p$ 
```

```
     $S \leftarrow S \cup \{p\}$ 
```

```
return  $S$ 
```

Implementation.

$O(n \log n)$ to sort.


$O(n)$ time for while loop.

Selecting Breakpoints: Correctness

Theorem. Greedy algorithm is optimal.

Pf. (by contradiction)

- Assume greedy is not optimal, and let's see what happens.
- Let $0 = g_0 < g_1 < \dots < g_p = L$ denote set of breakpoints chosen by greedy.
- Let $0 = f_0 < f_1 < \dots < f_q = L$ denote set of breakpoints in an optimal solution with $f_0 = g_0, f_1 = g_1, \dots, f_r = g_r$ for largest possible value of r .
- Note: $g_{r+1} > f_{r+1}$ by greedy choice of algorithm.

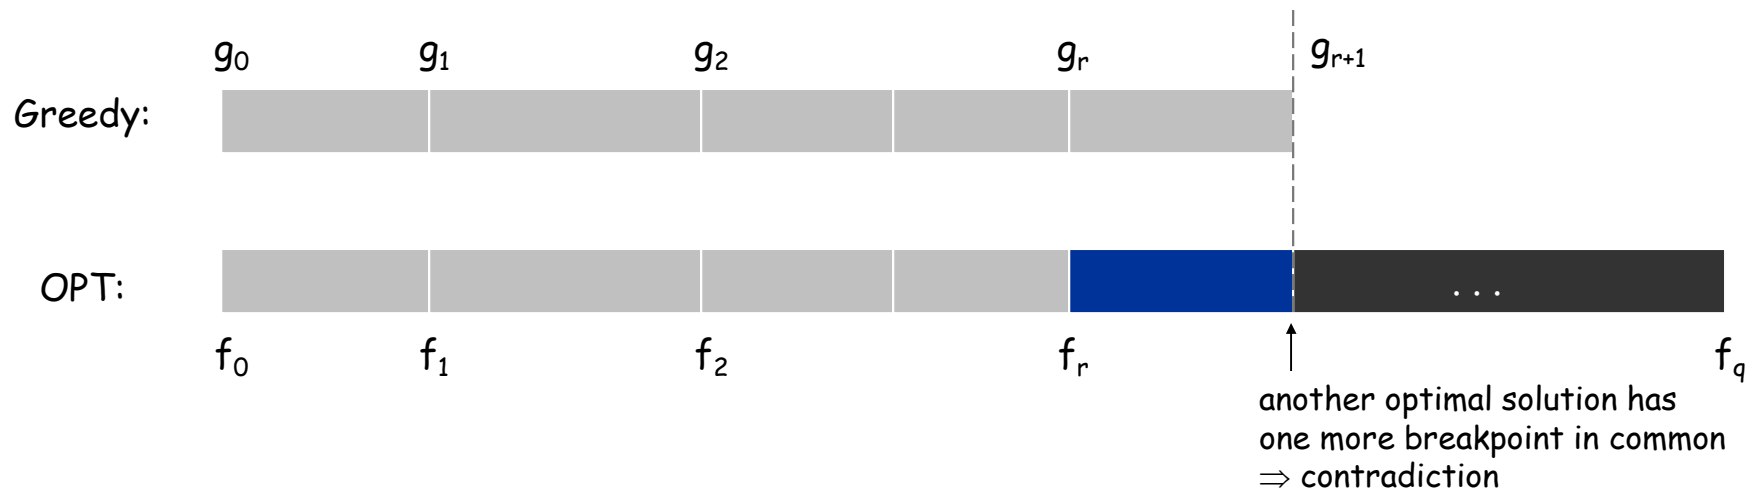


Selecting Breakpoints: Correctness

Theorem. Greedy algorithm is optimal.

Pf. (by contradiction)

- Assume greedy is not optimal, and let's see what happens.
- Let $0 = g_0 < g_1 < \dots < g_p = L$ denote set of breakpoints chosen by greedy.
- Let $0 = f_0 < f_1 < \dots < f_q = L$ denote set of breakpoints in an optimal solution with $f_0 = g_0, f_1 = g_1, \dots, f_r = g_r$ for largest possible value of r .
- Note: $g_{r+1} > f_{r+1}$ by greedy choice of algorithm.



Coin Changing

Goal. Given currency denominations: 1, 5, 10, 25, 100 (penny, nickel, dime, quarter, dollar) devise a method to pay amount to customer using fewest number of coins.

Ex: 34¢.



Cashier's algorithm. At each iteration, add coin of the largest value that does not take us past the amount to be paid.

Ex: \$2.89.



Coin-Changing: Greedy Algorithm

Cashier's algorithm. At each iteration, add coin of the largest value that does not take us past the amount to be paid.

```
Sort coins denominations by value:  $c_1 < c_2 < \dots < c_n$ .  
  
    coins selected  
    ↙  
S  $\leftarrow \emptyset$   
while ( $x \neq 0$ ) {  
    let k be largest integer such that  $c_k \leq x$   
    if ( $k = 0$ )  
        return "no solution found"  
     $x \leftarrow x - c_k$   
     $S \leftarrow S \cup \{k\}$   
}  
return S
```

Q. Is cashier's algorithm optimal?

Coin-Changing: Analysis of Greedy Algorithm

Theorem. Greedy is optimal for U.S. coinage: 1, 5, 10, 25, 100.

Pf. (by induction on x)

- Consider optimal way to change $c_k \leq x < c_{k+1}$: greedy takes coin k .
- We claim that any optimal solution must also take coin k .
 - if not, it needs enough coins of type c_1, \dots, c_{k-1} to add up to x
 - table below indicates no optimal solution can do this
- Problem reduces to coin-changing $x - c_k$ cents, which, by induction, is optimally solved by greedy algorithm. ▪

k	c_k	All optimal solutions must satisfy	Max value of coins 1, 2, ..., $k-1$ in any OPT
1	1	$P \leq 4$	-
2	5	$N \leq 1$	4
3	10	$N + D \leq 2$	$4 + 5 = 9$
4	25	$Q \leq 3$	$20 + 4 = 24$
5	100	no limit	$75 + 24 = 99$

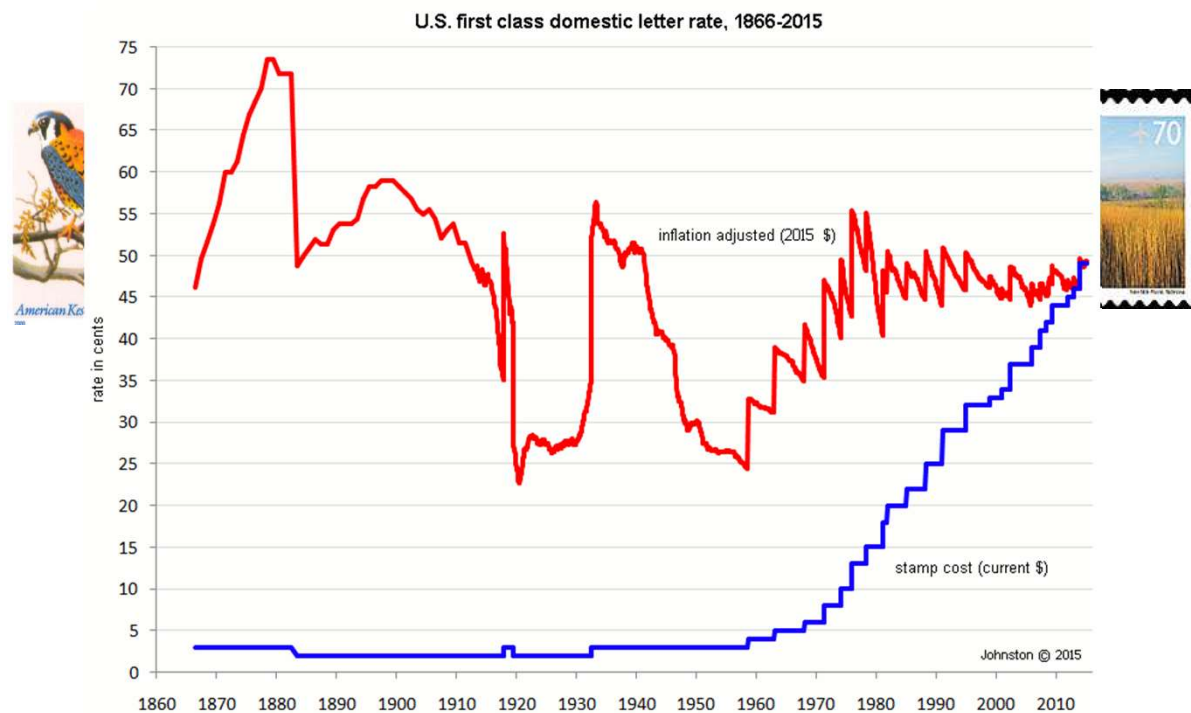
- Opt must have $p \leq 4$, because if $p \geq 5$, can replace 5 p with 1 n .
- Opt must have $n \leq 1$, because if $n \geq 2$, can replace with one dime, etc.
- So if don't use c_k , then must use c_1, \dots, c_{k-1} , and these must add up to $\geq x \geq c_k$.
- By case analysis, we see coins of type c_1, \dots, c_{k-1} never add up to $\geq c_k$.

Coin-Changing: Analysis of Greedy Algorithm

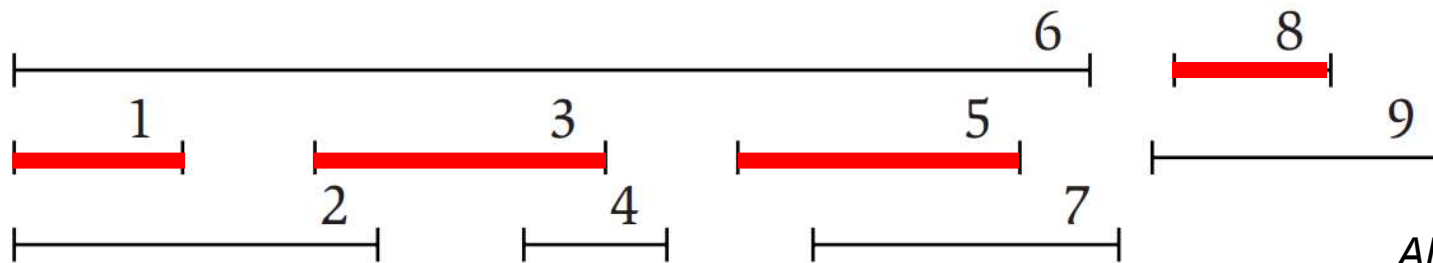
Observation. Greedy algorithm is sub-optimal for US postal denominations: 1, 10, 21, 34, 70, 100, 350, 1225, 1500.

Counterexample. 140¢.

- Greedy: 100, 34, 1, 1, 1, 1, 1.
- Optimal: 70, 70.



Interval scheduling



Algorithm Design
Kleinberg, Tardos

- Given a set of intervals, pick the largest number of nonoverlapping ones.
 - Each interval given by a start and finishing time.
- Models use of a shared resource.
 - Ex 9 people want to use a room. Different people want to use it at different times. Let max number of people use room.



A greedy algorithm

- Let's pick the intervals from left to right.
- **Intuition** Since can't pick a new interval until the previous one ends, want to pick intervals that end as quickly as possible.
- So we sort the intervals by finishing times. Then keep selecting **earliest finishing one** that doesn't overlap previous selected interval.

Initially let R be the set of all requests, and let A be empty

While R is not yet empty

 Choose a request $i \in R$ that has the smallest finishing time

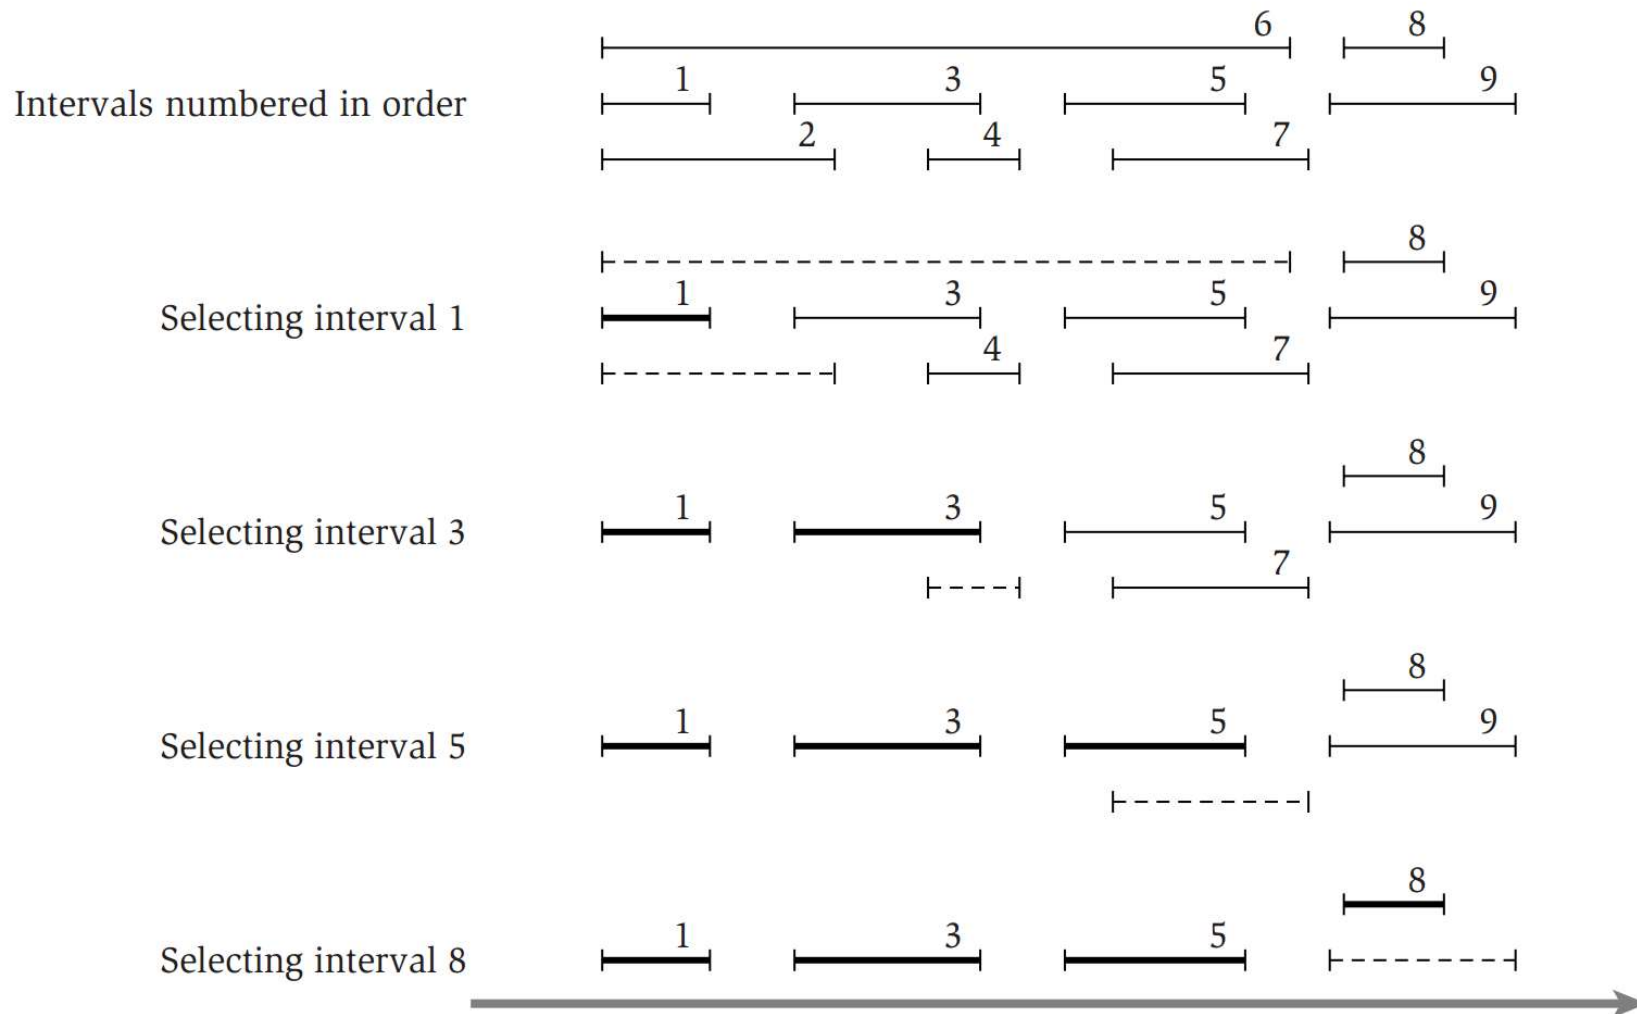
 Add request i to A

 Delete all requests from R that are not compatible with request i

EndWhile


Return the set A as the set of accepted requests

A greedy algorithm



Correctness

- We'll compare algorithm's solution S to an optimal solution T .
 - Can have $S \neq T$, since there may be several optimal solutions.
- **Def** Let s_k and t_k be k 'th interval in S and T , resp.
- **Def** let $\text{fin}(i)$ be finishing time of an interval i .
- **Claim** $\text{fin}(s_k) \leq \text{fin}(t_k)$ for all k .
- **Proof** By induction. True for $k=1$ since s_k is interval with min fin time.
 - Suppose true for $< k$, i.e. $\text{fin}(s_{k-1}) \leq \text{fin}(t_{k-1})$.
 - Then $\{\text{intervals not intersecting } s_{k-1} \text{ and finishing after } s_{k-1}\} \supseteq \{\text{intervals not intersecting } t_{k-1} \text{ and finishing after } t_{k-1}\}$.
 - By the algorithm, s_k is earliest finishing interval in the first set.
 - t_k is some interval in the latter set.
 - So $\text{fin}(s_k) \leq \text{fin}(t_k)$.
- **Corollary** S has at least as many intervals as T , i.e. S is optimal.



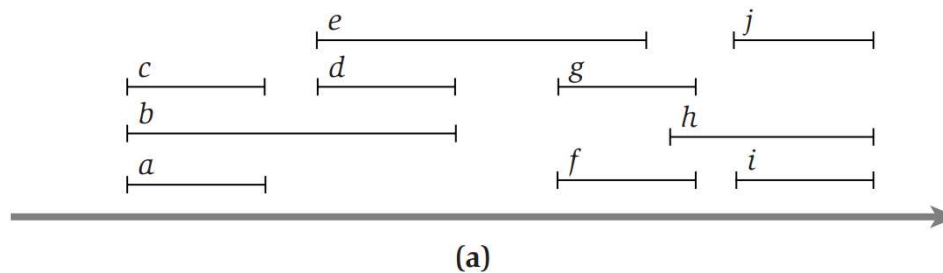


Analysis

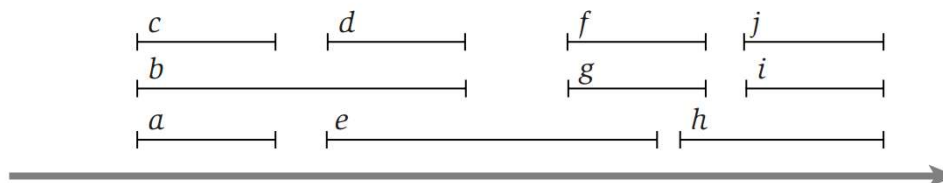
- If there are n intervals, sorting the intervals takes $O(n \log n)$ time.
- Then we go through the intervals in order of finishing times.
 - For each interval, check if it intersects last selected one, and select it if it doesn't.
 - Takes $O(n)$ time.
- Total $O(n \log n)$ time.

Interval coloring problem

- In interval scheduling, we scheduled the max number of intervals on one resource.
- In interval coloring, we need to schedule **all the intervals** on some number of resources.
 - Intervals on the same resource cannot overlap.
 - **Goal** Minimize the number of resources used.
- Example application is to schedule people who need to use a room in the min number of rooms.



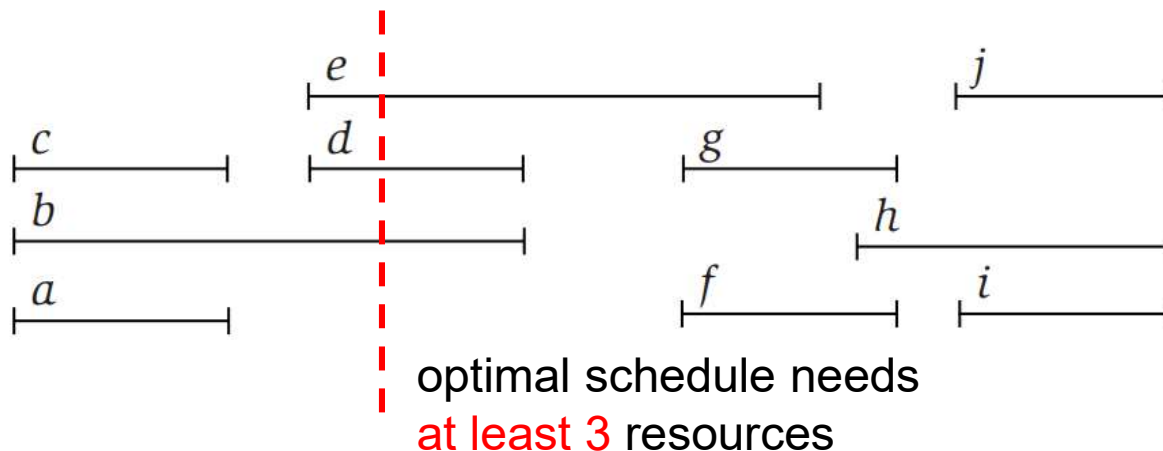
All intervals scheduled using **4 resources**



All intervals scheduled using **3 resources**

Optimality criterion

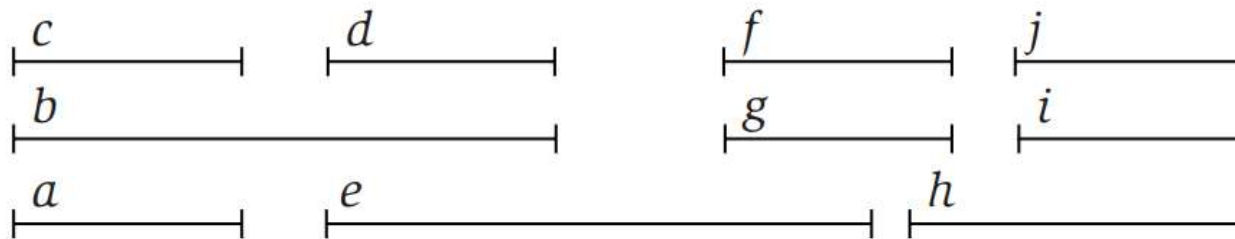
- **Observation** Suppose k intervals intersect at some time point. Then the optimal schedule needs **at least k** resources.
 - Since intervals on same resource can't overlap, then the k intersecting intervals need to be assigned to k different resources in any solution.
- **Def Depth** of a set of intervals is the max number of intervals that intersect at any time.



These intervals
have depth 3

Optimality criterion


- **Corollary** Let d be the depth of a set of intervals, and suppose we find a schedule using d resources. Then the schedule is optimal.
 - By the observation, any schedule needs at least d resources. Since our schedule uses d resources, it's optimal.



Since depth is 3,
this schedule is
optimal

A greedy algorithm

- Sweep through intervals in order of increasing start time.
 - Break ties arbitrarily.
- For each interval, assign it to smallest resource not already assigned to an intersecting interval.





Correctness

- **Claim** Let the set of intervals have depth d . Then the algorithm uses d resources.
- **Proof** Suppose algorithm is processing some interval s . Then $\leq d-1$ intervals intersect s .
 - So $\leq d-1$ resources assigned to these intervals.
 - So s can be assigned some resource $\leq d$.
- **Corollary** The algorithm is optimal
 - Follows by the optimality criterion.

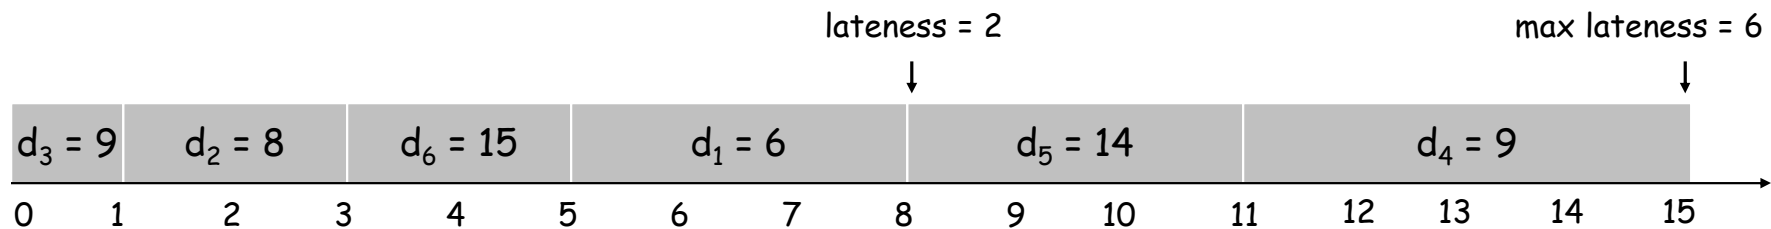
Scheduling to Minimizing Lateness

Minimizing lateness problem.

- Single resource processes one job at a time.
- Job j requires t_j units of processing time and is due at time d_j .
- If j starts at time s_j , it finishes at time $f_j = s_j + t_j$.
- Lateness: $\ell_j = \max \{ 0, f_j - d_j \}$.
- Goal: schedule all jobs to minimize **maximum** lateness $L = \max \ell_j$.

Ex:

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15



Minimizing Lateness: Greedy Algorithms

Greedy template. Consider jobs in some order.

- [Shortest processing time first] Consider jobs in ascending order of processing time t_j .
- [Earliest deadline first] Consider jobs in ascending order of deadline d_j .
- [Smallest slack] Consider jobs in ascending order of slack $d_j - t_j$.

Minimizing Lateness: Greedy Algorithms

Greedy template. Consider jobs in some order.

- [Shortest processing time first] Consider jobs in ascending order of processing time t_j .

	1	2	
t_j	1	10	counterexample
d_j	100	10	

- [Smallest slack] Consider jobs in ascending order of slack $d_j - t_j$.

	1	2	
t_j	1	10	counterexample
d_j	2	10	

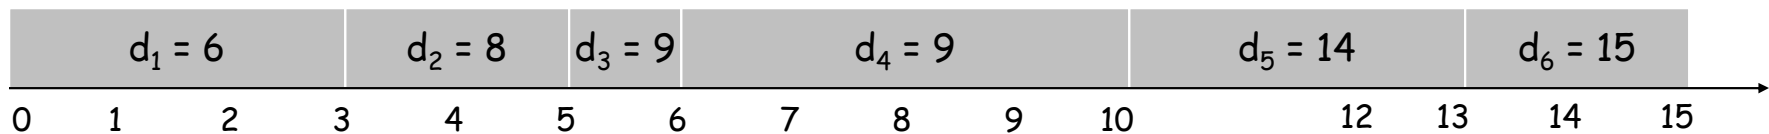
Minimizing Lateness: Greedy Algorithm

Greedy algorithm. Earliest deadline first.

```
Sort n jobs by deadline so that  $d_1 \leq d_2 \leq \dots \leq d_n$   
  
 $t \leftarrow 0$   
for  $j = 1$  to  $n$   
    Assign job  $j$  to interval  $[t, t + t_j]$   
     $s_j \leftarrow t, f_j \leftarrow t + t_j$   
     $t \leftarrow t + t_j$   
output intervals  $[s_j, f_j]$ 
```

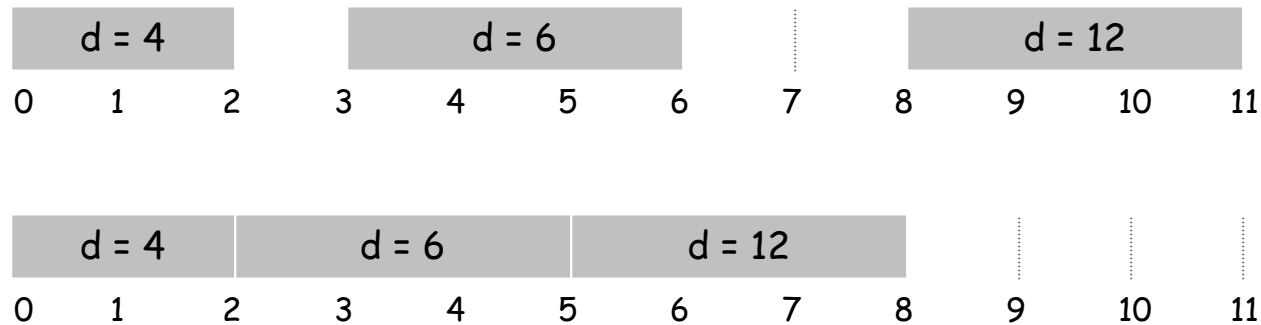
	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

max lateness = 1



Minimizing Lateness: No Idle Time

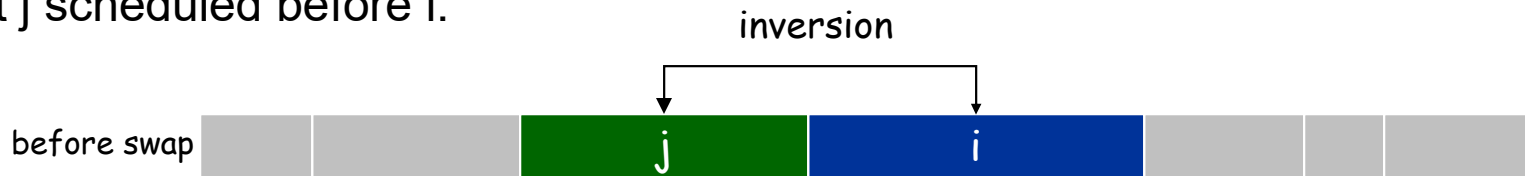
Observation. There exists an optimal schedule with no **idle time**.



Observation. The greedy schedule has no idle time.

Minimizing Lateness: Inversions

Def. An **inversion** in schedule S is a pair of jobs i and j such that: $i < j$ but j scheduled before i .

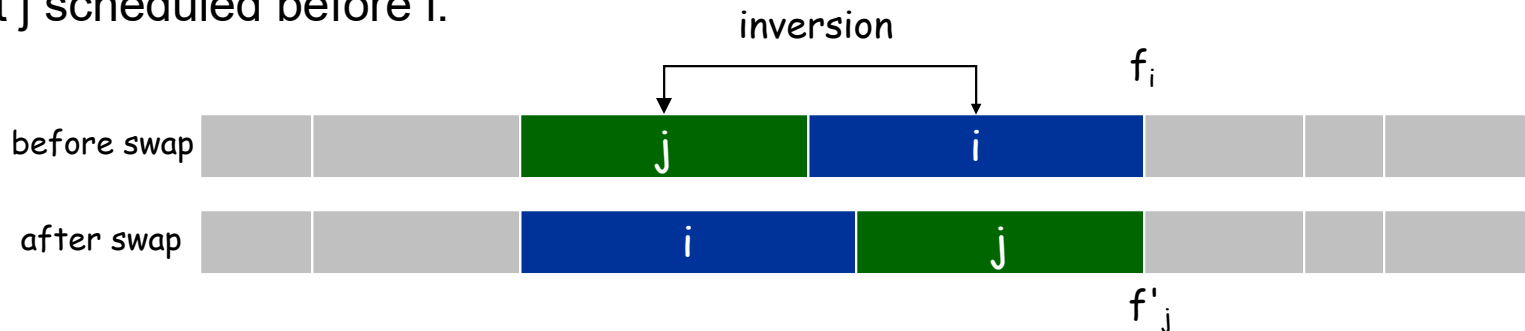


Observation. Greedy schedule has no inversions.

Observation. If a schedule (with no idle time) has an inversion, it has one with a pair of inverted jobs scheduled consecutively.

Minimizing Lateness: Inversions

Def. An **inversion** in schedule S is a pair of jobs i and j such that: $i < j$ but j scheduled before i .



Claim. Swapping two adjacent, inverted jobs reduces the number of inversions by one and does not increase the max lateness.

Pf. Let ℓ be the lateness before the swap, and let ℓ' be it afterwards.

- $\ell'_k = \ell_k$ for all $k \neq i, j$
- $\ell'_i \leq \ell_i$
- If job j is late:

$$\begin{aligned}
 \ell'_j &= f'_j - d_j && \text{(definition)} \\
 &= f_i - d_j && (j \text{ finishes at time } f_i) \\
 &\leq f_i - d_i && (i < j) \\
 &\leq \ell_i && \text{(definition)}
 \end{aligned}$$

Minimizing Lateness: Analysis of Greedy Algorithm

Theorem. Greedy schedule S is optimal.

Pf. Define S^* to be an optimal schedule that has the fewest number of inversions, and let's see what happens.

- Can assume S^* has no idle time.
- If S^* has no inversions, then $S = S^*$.
- If S^* has an inversion, let i - j be an adjacent inversion.
 - swapping i and j does not increase the maximum lateness and strictly decreases the number of inversions
 - this contradicts definition of S^* ▪