

Approximation algorithms 2

TSP, k-Center, Knapsack


CS240

Spring 2024

Rui Fan

Traveling Salesman Problem


- **Input** A complete graph with weights on the edges.
- **Output** A cycle that visits each city once.
- **Goal** Find a cycle with minimum total weight.



source: CLRS

Metric TSP

- TSP is NP-hard. In fact, it's even NP-hard to approximate when weights can be arbitrary.
- However, TSP is approximable for special types of weights.
- A weighted graph satisfies the triangle inequality if for any 3 vertices p, q, r , we have $d_{pq} + d_{qr} \geq d_{pr}$.
 - I.e., direct path is always no worse than a roundabout path.
 - This is called a metric TSP.
- There is a 1.5-approx algorithm for TSP in graphs with the triangle inequality.
 - Let's look at a simpler 2-approx first.










A 2-approximation for TSP

- Construct a minimum spanning tree T on G .
- Use depth-first traversal to visit all the vertices in T , starting from an arbitrary vertex.
- Convert this depth-first traversal T' to a cycle H that doesn't revisit any vertex.
- Return H as the TSP tour.

Example



source: CLRS



- (b) The MST T.
- (c) visit T in order abcbhbadefegeda.
- (d) converts the tour from (c) to a Hamiltonian cycle, that doesn't revisit any vertices.
- (e) is the optimal TSP.

Making the tour Hamiltonian

- To go from (c) to (d), we need to make a tour T' that revisits vertices into a cycle H that doesn't revisit vertices.
- We use shortcutting.
 - If we revisit a vertex in T' , we directly jump to the next vertex in T' we haven't visited.
 - We allow revisiting the first vertex.
 - The sequence of vertices we now visit is H .
 - Ex $abc**h**ba**de**fed**a** \rightarrow abchdefga$.





Making the tour Hamiltonian

- **Lemma** If H is the shortcut of T' , then $c(H) \leq c(T')$.
- **Proof** We formed H from T' by skipping over some vertices. E.g. we directly went from c to h , skipping over b .
 - But by the triangle inequality, $d_{cb} + d_{bh} \geq d_{ch}$.
 - So shortcutting from c to h didn't increase the distance.
 - The same thing applies to all our shortcuts.
 - So H is no longer than T' .





Proof of 2-approximation

- Let H^* be an optimum TSP.
- If we delete an edge from H^* , we get a spanning tree.
- Since T is an MST, $c(T) \leq c(H^*)$.
- Call the path from the depth-first traversal T' .
 - T' crosses each edge in T twice.
 - So $c(T') = 2 c(T)$.
- Let H be the outcome of shortcutting T' .
 - H is a Hamiltonian cycle. It visits all the vertices, and ends where it started.
 - $c(H) \leq c(T')$, by the lemma.
 - $c(H) \leq c(T') = 2 c(T) \leq 2 c(H^*)$.
- So H is a 2-approximation.



Matchings and Euler cycles

- A matching in a graph is a set of nonintersecting edges.
 - A perfect matching is a matching that includes every vertex.
- An Euler tour of a graph is a path that starts and ends at the same vertex, and visits every edge once.
 - Hamiltonian tour visits every vertex once.
- **Thm** (Euler) A graph has an Euler tour if and only if all vertices have even degree.
- Note how deciding if graph has Euler tour is trivial, but deciding if it has Hamiltonian tour is NPC!




Christofides 3/2-approx algorithm

- ❖ A 3/2-approximation for TSP with triangle inequality.
- Construct a minimum spanning tree T on G .
- Find the set V' of odd degree vertices in T .
- Construct a minimum cost perfect matching M on V' .
- Add M to T to obtain T' .
- Find an Euler tour T'' in T' .
- Shortcut T'' to obtain a Hamiltonian cycle H . Output as the TSP.



Why Christofides works well

- In the 2-approx, we found a TSP by “doubling” the MST to an Euler tour, then shortcutting.
 - We need to start with Euler tour before shortcutting to ensure we visit all cities.



- Key to Christofides is to find a shorter Euler tour, without doubling the MST.
 - A graph with only even degree vertices always has Euler tour.
 - So we want to modify the MST to have all even degrees, by adding a matching.




Proof of correctness

- **Lemma** T' has an Euler tour.
- **Proof** There are an even number of vertices in V' , because the total degree of T is even.
 - Since G is a complete graph and $|V'|$ is even, there's a perfect matching on V' .
 - The min cost perfect matching can be found in $O(n^2)$ time using the blossom algorithm.
 - The degree of every node in M is odd. Since V' are the odd degree nodes in T , adding M to T makes all nodes in T' have even degree.
 - T' has Euler tour by Euler's theorem.

Proof of correctness

- **Lemma** Let H^* be an optimal TSP on G , and let m be the cost of M . Then $m \leq c(H^*)/2$.
- **Proof** Let H' be the optimal TSP on V' .
 - $c(H') \leq c(H^*)$ because H' is an optimal TSP on fewer vertices.
 - H' is a cycle on V' , so it consists of two matchings on V' . The cheaper one has cost $m' \leq c(H')/2 \leq c(H^*)/2$.
 - $m \leq m'$ because M has min cost.



Proof of 3/2-approximation

- **Thm** Let H be the TSP output by Christofides and let H^* be an optimal TSP. Then $c(H) \leq 3/2 \cdot c(H^*)$.



- **Proof**

- $c(T) \leq c(H^*)$ because T is an MST.
- $c(T') = c(M) + c(T) \leq c(H^*)/2 + c(H^*) = 3/2 \cdot c(H^*)$.
- $c(H) \leq c(T')$ because H is the shortcut of T' .

- Construct a minimum spanning tree T on G .
- Find set V' of odd-degree vertices in T .
- Construct a minimum cost perfect matching M on V' .
- Add M to T to obtain T' .
- Shortcut T' to obtain a Hamiltonian cycle. Output as the TSP.

k-Center problem

- Given a city with n sites, we want to build k centers to serve them.
 - Let S be set of sites, C be set of centers.
- Each site uses the center closest to it.
 - Distance of site s from the nearest center is $d(s, C) = \min_{c \in C} d(s, c)$.
- Goal is to make sure no site is too far from its center.
 - We want to minimize the max distance that any site is from its closest center.
 - Minimize $r(C) = \max_{s \in S} \min_{c \in C} d(s, c)$.
 - C is called a cover of S , and r is called C 's radius.
 - Where should we put centers to minimize the radius?
- Assume distances satisfy triangle inequality.



Gonzalez's algorithm

- k-Center is NP-complete.
- We'll give a simple 2-approximation for it.
- **Idea** Say there's one site that's farthest away from all centers. Then it makes the radius large. We'll put a center at that site, to reduce the radius.
 - Note we allow putting center at same location as site.






Gonzalez's algorithm

- C is set of centers, initially empty.
- repeat k times
 - choose site s with maximum $d(s, C)$
 - add s to C
- return C
- **Note** The centers are located at the sites.

Proof of correctness

- Let C be the algorithm's output, and r be C 's radius.
 - $r = \max_{s \in S} \min_{c \in C} d(s, c)$
- **Lemma 1** For any $c, c' \in C$, $d(c, c') \geq r$.
- **Proof** Since r is the radius, there exists a point $s \in S$ at distance $\geq r$ from all the centers.
 - If there's no such s , then C 's radius $< r$.
 - So s is distance $\geq r$ from c and c' .
 - Suppose WLOG c' is added to C after c .
 - If $d(c, c') < r$, then algorithm would add s to C instead of c' , since s is farther.






Proof of correctness

- **Cor** There exist $k+1$ points mutually at distance $\geq r$ from each other.
 - By the lemma, the k centers are mutually $\geq r$ distance apart.
 - Also, there's an $s \in S$ at distance $\geq r$ from all the centers.
 - Otherwise C 's covering radius is $< r$.
 - So the k centers plus s are the $k+1$ points.
- Call these $k+1$ points D .

Proof of correctness

- Let C^* be an optimal cover with radius r^* .
- **Lemma 2** Suppose $r > 2r^*$. Then for every $c \in D$, there exists a corresponding $c' \in C^*$. Furthermore, all these c' are unique.
- **Proof** Draw a circle of radius $r/2$ around each $c \in D$.
 - There must be a $c' \in C^*$ inside the circle, because
 - c is at most distance r^* away from its nearest center, since r^* is C^* 's radius.
 - $r/2 > r^*$.
 - Given $c_1, c_2 \in D$, let $c'_1, c'_2 \in C^*$ be inside c_1 and c_2 's circle, resp.
 - c_1 and c_2 's circles don't touch, because $d(c_1, c_2) \geq r$.
 - So $c'_1 \neq c'_2$.






Proof of correctness

- **Thm** Let C be the output of Gonzalez's algorithm, and let C^* be an optimal k -center. Then $r(C) \leq 2r(C^*)$.
- **Proof** By Lemma 2, if $r(C) > 2r(C^*)$, then for every $c \in D$, there is a unique $c' \in C^*$.
 - But there are $k+1$ points in D , by the corollary.
 - So there are $k+1$ points in C^* . This is a contradiction because C^* is a k -center.

The knapsack problem

- We have a set of items, each having a weight and a value.
- We have a knapsack that can carry up to W amount of weight.
- We want to put items in the knapsack to maximize the total value, but not exceed the weight limit.
- **Ex** Items 3 and 4 are the highest value items with weight ≤ 11 .
- Assume all items have weight $\leq W$, i.e. any single item fits in knapsack.



$W = 11$

Item	Value	Weight
1	1	1
2	6	2
3	18	5
4	22	6
5	28	7

A dynamic program for knapsack

- Let $\text{OPT}(i,v)$ = minimum weight of a subset of items $1,\dots,i$ that has value $\geq v$.
- If optimal solution uses item i .
 - Then we pay w_i weight for item i , and need to achieve value $\geq v-v_i$ using items $1,\dots,i-1$ using min weight.
 - So $\text{OPT}(i,v)=w_i+\text{OPT}(i-1,v-v_i)$.
- If optimal solution doesn't use item i .
 - Then we need to achieve value $\geq v$ using items $1,\dots,i-1$.
 - So $\text{OPT}(i,v)=\text{OPT}(i-1,v)$.
- Choose the case that gives smaller weight.
- $$\text{OPT}(i,v) = \begin{array}{ll} 0 & \text{if } v=0 \\ \infty & \text{if } i=0, v>0 \\ \min(\text{OPT}(i-1,v), w_i+\text{OPT}(i-1,v-v_i)) & \text{otherwise} \end{array}$$



Running time of dynamic program

- Say there are n items, and the largest value of any item is v^* .
- The max value we can pack into the knapsack is nv^* , where v^* is the largest v value.
- Solve all subproblems of the form $\text{OPT}(i,v)$, where $i \leq n$ and $v \leq nv^*$.
 - This is a total of $O(n^2v^*)$ subproblems.
- The solution to Knapsack is the max value V that can be packed with weight $\leq W$.
- Having solved all the subproblems, we can find V by finding the subproblem with the largest value that has optimum weight $\leq W$.
 - $V = \max_{v \leq nv^*} \text{OPT}(n,v) \leq W$.
- So solving Knapsack takes total time $O(n^2v^*)$.



Running time of dynamic program

- The DP gives an optimal solution to Knapsack and takes $O(n^2v^*)$ time. Have we found a polytime algorithm for an NP-complete problem?
- No. The problem size is $O(n \log(v^*))$, because it takes $\log(v^*)$ bits to express each item's value. But $O(n^2v^*)$ is not polynomial in $n \log(v^*)$.
- To make this DP fast, we have to make the largest value small.



PTAS

- Let $\varepsilon > 0$ be any number. We'll give a $(1+\varepsilon)$ -approximation for knapsack.
- By setting ε sufficiently small, we can get as good an approximation as we want!
 - This type of algorithm is called a polynomial time approximation scheme, or PTAS.
- Contrast this with earlier algs we studied, which had worse approx ratios, e.g. 2 or $\log n$.
- But the running time will be $O(n^3/\varepsilon)$. Hence we can't set $\varepsilon=0$ get the optimal solution.
- We're trading accuracy for time. The more accurate (smaller ε), the more time the algorithm takes.

Main idea: rounding

- Since we only need an approximate solution, we can change the values of the items a little (round the values) and not affect the solution much.
- We scale and round the original values to make them small.
- The previous DP took $O(n^2v^*)$ time. So if the rounded values are small, this DP is fast.

W = 11						W = 11		
Item	Value	Weight				Item	Value	Weight
1	134,221	1				1	2	1
2	656,342	2				2	7	2
3	1,810,013	5				3	19	5
4	22,217,800	6				4	223	6
5	28,343,199	7				5	284	7

Rounding

- Let $\varepsilon > 0$ be the precision we want.
- Set $\theta = \frac{\varepsilon v^*}{2n}$ to be a scaling factor.
 - v^* is the largest value of any item.
- Scale all values down by θ then round up.
 - $v' = \lceil v/\theta \rceil$.
- Make a problem where each value v_i is replaced by v'_i .
 - Call this the scaled rounded problem.
- Let v^\wedge be max value in the scaled rounded problem. Then $v^\wedge = \lceil v^*/\theta \rceil = \lceil v^*/(\varepsilon v^*/2n) \rceil = \lceil 2n/\varepsilon \rceil$.
- Running time of DP on scaled rounded problem is $O(n^2 v^\wedge) = O(n^3/\varepsilon)$.



Solving the original problem

- Make another new problem in which each value v_i is replaced by $u_i = \lceil v_i/\theta \rceil * \theta$.
 - Call this the rounded problem.
 - We have $u_i \geq v_i$, and $u_i \leq v_i + \theta$.
- Note u values are equal to v' values multiplied by θ .
 - Thus, the optimal solution for the rounded problem and the scaled rounded problem are the same.
- We now have 3 problems, the original problem, the scaled rounded problem, and the rounded problem.
- Let S be the optimal solution to the scaled rounded problem, which we can find in time $O(n^3/\varepsilon)$. S is also optimal for the rounded problem.
- We'll show S is a $1+\varepsilon$ approximation for the original problem.

Correctness

- **Thm** Let S^* be the optimal solution to the original problem. Then $(1+\varepsilon) \sum_{i \in S} v_i \geq \sum_{i \in S^*} v_i$.
Hence S is a $(1+\varepsilon)$ -approximate solution.

- **Proof**

$$\sum_{i \in S^*} v_i \leq \sum_{i \in S^*} u_i$$

$$u_i \geq v_i$$

$$\leq \sum_{i \in S} u_i$$

S is opt soln for rounded problem

$$\leq \sum_{i \in S} (v_i + \theta)$$

$$u_i \leq v_i + \theta$$

$$\leq \sum_{i \in S} v_i + n \theta$$

$$|S| \leq n$$

Correctness

- Suppose item j has the largest value, so

$$v^* = v_j. \text{ Then } n\theta = \frac{\varepsilon}{2} v_j \leq \frac{\varepsilon}{2} u_j \leq \frac{\varepsilon}{2} \sum_{i \in S} u_i$$

□ Last inequality because item j itself is feasible solution, so opt solution S is no smaller.

- So $\sum_{i \in S} v_i \geq \sum_{i \in S} u_i - n\theta \geq \left(\frac{2}{\varepsilon} - 1\right) n\theta$, where first inequality comes from inequalities on last page.

- Assuming $\varepsilon \leq 1$, then $n\theta \leq \varepsilon \sum_{i \in S} v_i$

- Finally, we have

$$\sum_{i \in S^*} v_i \leq \sum_{i \in S} v_i + n\theta \leq \sum_{i \in S} v_i + \varepsilon \sum_{i \in S} v_i = (1 + \varepsilon) \sum_{i \in S} v_i$$



Summary

- We gave a DP for Knapsack.
- We scale and round to reduce number of different item values.
- Running the DP on the scaled rounded problem and using the solution for the original problem leads to an arbitrarily good approximation for Knapsack, a PTAS.
- There are PTAS's for a number of other problems.
 - Multiprocessor scheduling.
 - Bin packing.
 - Euclidean TSP.
- However, there are also many problems for which PTAS's do not exist, unless $P=NP$.